

Report 1

Introduction

Transposing

OMP

Blocking

Unrolling

FMA: NEON, AVX (intel/apple)

Cache Locality with ikj

Apple Accelerate

Macbook pro M3

L1 instruction cache: 192KB (hw.perflevel0.l1icachesize: 196608)

L1 data cache: 128KB (hw.perflevel0.l1dcachesize: 131072)

L2 cache: 16MB (hw.perflevel0.l2cachesize: 16777216)

How we validate

By using numpy to write the output of `A @ B` to `tmp/matmul` we can read the and compare the output of our GEMM

How we measure improvements

Gflops, our base metrics we will comparing with is Numpy with OpenBlas, however Numpy has the capability to do gemm with CBlas and Apple Accelerate framework. We will have some comparisns with OpenBlas and Cblas, but the main Numpy "framework" will be OpenBlas.

Run commands and flags?

```
-std=c++17 -Xpreprocessor -fopenmp -O3 -march=native -mcpu=native -ffast-math -I/opt/homebrew/include -L/opt/homebrew/lib -lomp -o
```

Naive implementation

Python

```
for i in range(N):
    for j in range(N):
        for k in range(N):
            C[i][j] += A[i][k] * B[k][j]
```

average Gflops: 0.015153

N: 256

cpp

```
void gemm_stupid(){
    for (int i = 0; i < N; i++)
        for (int j = 0; j < N; j++) {
            for (int k = 0; k < N; k++)
                C[i][j] += A[i][k] * B[k][j];
        }
}
```

average Gflops: 3.738789

N: 256

Numpy OpenBlas (single threaded)

To confirm that Numpy is using OpenBlas we can view the output of `print(np.__config__.show())` and we see the output include name: `scipy-openblas`. depending on how we installed numpy to the venv we would sometime get `Accelerate` or `cblas`.

We use the library `threadpoolctl` to force numpy into only using a single thread as shown below.

```
from threadpoolctl import threadpool_limits

with threadpool_limits(limits=1):
    pass
```

We use numpys `@` operator to do the matrix multiplication.

It is also important to note that these matrices are initialized as type `float32` as shown below. This is to give numpy a fair chance since we are using float in Cpp

```
A = np.random.randn(N, N).astype(np.float32)
B = np.random.randn(N, N).astype(np.float32)

with threadpool_limits(limits=1):
    for i in range(itr):

        start = monotonic()
        C = A @ B
        end = monotonic()
        s = end - start

        gflops = (N*N*2*N) / (s * 10**9)
        avg += gflops

    print(f"Gflops: {gflops:.6f}")

    print(f"avg: {avg/itr:.2f}")
```

average Gflops: 43.073460

N: 2048

Numpy (multi threaded)


Numpy by default is multithreaded so we will be testing out how good default numpy is by removing

```
with threadpool_limits(limits=1):
```

We can also confirm that this is multithreaded by checking `htop`

This gives us 173.38 avg Gflops with N = 2048

Lets also confirm that this is multi threaded.



Here is the current table

Descriptiopn	Gflops	N
Python Stupid	0.015	256
C++ Stupid	3.738	256
ST_Numpy	43.073	2048
Numpy	173.38	2048

Improving CPP: Transposing

By transposing **B**, the data we fetch should be more continuous, thus reducing our cache misses and being generally more cache friendly.

Having the matrices 1D should also help with cache misses, but with our testing it does not give us any preformance advantage, rather it makes it a bit slower. So we will stick with 2D arrays.

We think that the -O2 / -O3 might be optimizing the 2D arrays better.

We get an average of 35.651 Gflops for the transposed version, which is a huge improvment over the non transpose.

Descriptiopn	Gflops	N
Python Stupid	0.015	256
C++ Stupid	3.738	256
Numpy ST	43.073	2048
Numpy MT	173.38	2048
C++ Transposed	35.651	2048

Improving CPP: Blocking (Loop tiling)

Next we are trying blocking, we will add three outer loops to create a smaller matrix that we focus on, then tiling up as we finish the matrix.

```
void gemm_omp() {
    int bi, bk, bj, i, k, j;

    for(bi = 0; bi < N; bi += BLOCKSIZE)
        for(bj = 0; bj < N; bj += BLOCKSIZE)
            for(bk = 0; bk < N; bk += BLOCKSIZE)

                for(i = 0; i < BLOCKSIZE; i++)
                    for(j = 0; j < BLOCKSIZE; j++) {
                        for(k = 0; k < BLOCKSIZE; k+=1) {
                            C[bi + i][bj + j] += A[bi + i][bk + k] * BT[bj + j][bk + k];
                        }
                    }
}
```

why does blocking even help? (images and memory)

on book page 268, 6.4 Hardware exploration

theoretical best based on cache size

testing best

blockSize	Matrix Size	FLOP/s
4	2048	31.93 GFLOPS
8	2048	24.00 GFLOPS
16	2048	31.15 GFLOPS
32	2048	25.36 GFLOPS
64	2048	35.90 GFLOPS
128	2048	33.29 GFLOPS
256	2048	24.18 GFLOPS
512	2048	23.18 GFLOPS

It is important to know that this is the minimum we can do with blocking, and there is a lot of room for improvement. For example, we can optimize our tiling based on the cache size, we could also tile X and Y differently, optimizing for L1 and L2 caches, however this is very complicated. We might look into this if we have more time.

GRAPH THE ABOVE

Threading with OMP EXPLAIN HERE MORE

Before going over to doing Blocking and Unrolling, we will test out the transpose version with OMP threading. We are doing this because we think that Blocking and Unrolling will have a much bigger affect once we do gemm with threading.

The simplest OMP we could use is `#pragma omp parallel for` on the outer most loop,

To do even better, we can

```
#pragma omp parallel for collapse(2) private(bj, bk, i, j) shared(A, BT, C)
```

We also did testing on blocksize with threading as it might not be the same.

blockSize	Matrix Size	FLOP/s
4	2048	198.258 GFLOPS
8	2048	132.660 GFLOPS
16	2048	60.810 GFLOPS
32	2048	62.349 GFLOPS
64	2048	169.871 GFLOPS
128	2048	224.167 GFLOPS
256	2048	121.976 GFLOPS

We see here that these numbers are much better than single threaded. It turns out that the cache advantage given by blocking is utilized much more when threaded.

Average flops: 224.167 with block size 128

As we observe here we are very close to numpy, even beating by a couple of Gflops. However, this is not because we are fast, rather it is because Numpy with OpenBlas is quite slow. We have beaten Numpy here but we will keep pushing for faster and faster.

Unrolling

Unrolling is a very nice idea that will help us remove the inner most loop by unrolling it.

- doing the if statement logic (compiler wise) is expensive

our current inner loop looks like this

```

..
for (j = 0; j < blockSize; j++)
    for (k = 0; k < blockSize; k++){
        C[bi + i][bj + j] += A[bi + i][bk + k] * B_trans[bj + j][bk + k];
    }
..

```

We have an issue where the inner most if loop is being updated for each entry of `k` and this is expensive since the compiler needs to do extra work, and extra memory usage is done here to fetch and update `k`

we can reduce this workload by unrolling `x` amount of instructions and writing them down manually.

for example unrolling by 4 instructions would look like this

```

..
for (j = 0; j < blockSize; j++)
    for (k = 0; k < blockSize; k+=4){
        C[bi + i][bj + j] += A[bi + i][bk + 0] * B_trans[bj + j][bk + 0];
        C[bi + i][bj + j] += A[bi + i][bk + 1] * B_trans[bj + j][bk + 1];
        C[bi + i][bj + j] += A[bi + i][bk + 2] * B_trans[bj + j][bk + 2];
        C[bi + i][bj + j] += A[bi + i][bk + 3] * B_trans[bj + j][bk + 3];
    }
..

```

However, our testing shows that this barely gives any performance advantage. Neither did unrolling 8 instructions.

The problem is that we are still using the if statement, so the best scenario would be to unroll all the instructions and remove the if statement.

the innermost loop runs up to blocksize, meaning there would be the same amount of instructions as the block size, so we need to unroll 16 instructions

to be able to remove the if statement

```

..
for (j = 0; j < blockSize; j++) {
    C[bi + i][bj + j] += A[bi + i][bk + k] * B_trans[bj + j][bk + k];
    C[bi + i][bj + j] += A[bi + i][bk + k + 1] * B_trans[bj + j][bk + k + 1];
    C[bi + i][bj + j] += A[bi + i][bk + k + 2] * B_trans[bj + j][bk + k + 2];
    C[bi + i][bj + j] += A[bi + i][bk + k + 3] * B_trans[bj + j][bk + k + 3];

    C[bi + i][bj + j] += A[bi + i][bk + k + 4] * B_trans[bj + j][bk + k + 4];
    C[bi + i][bj + j] += A[bi + i][bk + k + 5] * B_trans[bj + j][bk + k + 5];
    C[bi + i][bj + j] += A[bi + i][bk + k + 6] * B_trans[bj + j][bk + k + 6];
    C[bi + i][bj + j] += A[bi + i][bk + k + 7] * B_trans[bj + j][bk + k + 7];

    C[bi + i][bj + j] += A[bi + i][bk + k + 8] * B_trans[bj + j][bk + k + 8];
    C[bi + i][bj + j] += A[bi + i][bk + k + 9] * B_trans[bj + j][bk + k + 9];
    C[bi + i][bj + j] += A[bi + i][bk + k + 10] * B_trans[bj + j][bk + k + 10];
    C[bi + i][bj + j] += A[bi + i][bk + k + 11] * B_trans[bj + j][bk + k + 11];

    C[bi + i][bj + j] += A[bi + i][bk + k + 12] * B_trans[bj + j][bk + k + 12];
    C[bi + i][bj + j] += A[bi + i][bk + k + 13] * B_trans[bj + j][bk + k + 13];
    C[bi + i][bj + j] += A[bi + i][bk + k + 14] * B_trans[bj + j][bk + k + 14];
    C[bi + i][bj + j] += A[bi + i][bk + k + 15] * B_trans[bj + j][bk + k + 15];
}
..

```

Surprising, this also gave us very little performance advantage, so we played around with the block size and found that a blocksize of 4 (corresponding to 4 unrolled instructions) gives the best performance

Final loop:

```
..  
for (j = 0; j < blockSize; j++) {  
    // Fully unrolled for blockSize = 4  
    C[bi + i][bj + j] += A[bi + i][bk + 0] * B_trans[bj + j][bk + 0];  
    C[bi + i][bj + j] += A[bi + i][bk + 1] * B_trans[bj + j][bk + 1];  
    C[bi + i][bj + j] += A[bi + i][bk + 2] * B_trans[bj + j][bk + 2];  
    C[bi + i][bj + j] += A[bi + i][bk + 3] * B_trans[bj + j][bk + 3];  
}  
..
```

average gflops: 299.380005

N = 2048

Table update:

Descriptiogn	Gflops	N
Python Stupid	0.015	256
C++ Stupid	3.738	256
Numpy ST	43.073	2048
Numpy MT	173.38	2048
C++ Transposed	35.651	2048
C++ Transposed, OMP, Blocking=128	224.167	2048
C++ Transposed, OMP, Unrolled (blocking 4)	290.167	2048

FMA: NEON, AMX

Fused Multiply-add is an instruction that can execute $X \leftarrow ax + b$ in the same amount as doing addition and the multiplication, this is could potentially help the processor with asymptotic speed per clock cycle[1]

At first we were not sure if -O3 was enabling FMA operations by default, so we tried the C++ Transposed and unrolled version with the `-mno-fma` and got 291.007 Gflops, and with the `mfma` we got the same amount.

So the compiler might not be doing this by default.

we added `#pragma omp simd` to help the compiler and got a consistent 5-6 Gflop increase,

```
#pragma omp simd  
for (j = 0; j < blockSize; j++) {  
    C[bi + i][bj + j] += A[bi + i][bk + 0] * BT[bj + j][bk + 0];  
    C[bi + i][bj + j] += A[bi + i][bk + 1] * BT[bj + j][bk + 1];  
    C[bi + i][bj + j] += A[bi + i][bk + 2] * BT[bj + j][bk + 2];  
    C[bi + i][bj + j] += A[bi + i][bk + 3] * BT[bj + j][bk + 3];  
}
```


talk about below (page 78)

`#pragma omp simd` vs `#pragma vector` always

Looking at the assembly code on a cpp explorer (<https://godbolt.org/noscript>) we can find a single mention of `fmadd`, suggesting that the compiler is using fma, even though the speed up of 5-8 Gflops is underwhelming - maybe we are not using its full potential?

Cache locality

After further investigation we realised that we are calculating the matrix in the wrong way, initially we were using the ijk approach, but the ikj approach should have a much better cache locality

Users/matar/School/4DT906-Parallel-and-Distributed/Assignment1/report/images/ijk_ikj.png

EXPLAIN MORE WHY ITS GOOD'

ALSO UN-UNROLLOED; PRAGMA SMD FIXES ALREADY??

cant transpose anymore since loop is different

same thing? but better results


doing this gives an impressive boost up to 410.236 Gflop!

Descriptiopl	Gflops	N
Python Stupid	0.015	256
C++ Stupid	3.738	256
Numpy ST	43.073	2048
Numpy MT	173.38	2048
C++ Transposed	35.651	2048
C++ Transposed, OMP, Blocking=128	224.167	2048
C++ Transposed, OMP, Unrolled	290.167	2048
C++, OMP, cache locality	410.236	2048


Larger N and memory bound

REDO THIS GRAPH DUH

While we are very close to numpy with what we have, once we scale up N to 4096 and larger, we can see a quite large decrease in Flops.

Users/matar/School/4DT906-Parallel-and-Distributed/Assignment1/report/images/openblas.png

Note that the below graph is CBLAS and not OpenBlas, how are they so fast for large N?

Users/matar/School/4DT906-Parallel-and-Distributed/Assignment1/report/images/Figure_1.png

We think that this is due to being memory bound, or the block size = 4 being limiting. so lets try to use larger block size

Apple Accelerate framework